

Problem definition

The goal of the project was to create a RL Agent which is able to train from scratch by interacting with the Reacher (robot arm) environment for a series of episodes. During the episode the agent receives the environment state. The agent is able to perform some actions with may change the environment state.

Environment description

The **state space** has 37 dimensions each of which is a continuous variable. It includes the agent's velocity, along with ray-based perception of objects around the agent's forward direction. The **action space** contains the following 4 legal actions:

- move forward 0
- move backward 1
- turn left 2
- turn right 3

A reward of +1 is provided for collecting a yellow banana, and a reward of -1 is provided for collecting a blue banana. Thus, the goal of the agent is to collect as many yellow bananas as possible while avoiding blue bananas. The task is episodic, and in order to solve the environment, the agent must get an average score of +13 over 100 consecutive episodes.

Solution

Deep Q-Network Architecture

Network

Since the agent should act upon environment state provided as vector of size 37, I decided to build simple network with 3 fully connected layers, which should predict the Q-value for each action given the input state. The input has size 37, output 4. I tried 64x64, 64x32, 32x32, 32x16 hidden layers configuration. The best one was 32x32, since it was able to reach the training goal faster and it has less internal parameters then 64x64 and 64x32.

Training

Loss function and optimizing method

I tried 2 loss functions 1: Simple DQN loss 2 and Double DQN loss 3

$$L \equiv (Y_t - Q(S_t, a; \theta_t))^2 \quad (1)$$

$$Y_t \equiv R_{t+1} + \gamma \max_a Q(S_t, a; \theta'_t) \quad (2)$$

$$Y_t \equiv R_{t+1} + \gamma Q(S_{t+1}, \operatorname{argmax}_a Q(S_t, a; \theta_t); \theta'_t) \quad (3)$$

Where $Q(S_t, a, \theta_t)$ is output (value of action a) of the network with internal parameters θ_t given input S_t . Both loss functions used 2 sets of network parameters: target θ'_t and local θ_t . Target network parameters was updating through soft update $\theta'_t = (1 - \tau)\theta'_t + \tau\theta_t$. As optimizing method for θ_t I used Adam with learning rate α .

Experience replay

The input for training did not come directly from episodes. Instead I used experience replay technique (see [Human-level control through deep reinforcement learning](#)).

Results

I came up using following training hyper parameters:

- ϵ_{start} : 1.0, ϵ_{end} : 0.01, ϵ_{decay} : 0.97
- Buffer size: 100000
- Batch size: 64
- γ : 0.99
- τ : 0.005
- α : 0.0005

They were used across all training experiments below.

Training DDQN 32x16

DQN architecture: fully connected with bias 37x32x16x4. Loss [3](#).

Training history:

Episode 100. Average Score: 4.88. Time elapsed: 4:35

Episode 200. Average Score: 12.21. Time elapsed: 9:11

Environment solved in 170 episodes! Average Score: 13.12. Time elapsed: 12:31s. See [1](#).




Figure 1: DDQN 32x16: Rewards per episode

Training DDQN 32x32

DQN architecture: fully connected with bias 37x32x32x4. Loss [3](#).

Training history:

Episode 100. Average Score: 6.50. Time elapsed: 3:26

Episode 200. Average Score: 12.66. Time elapsed: 6:51

Environment solved in 109 episodes! Average Score: 13.08. Time elapsed: 7:09s. See [2](#).

Training DDQN 64x32

DQN architecture: fully connected with bias 37x64x32x4. Loss [3](#).

Training history:

Episode 100. Average Score: 5.69. Time elapsed: 3:25

Episode 200. Average Score: 12.38. Time elapsed: 6:52

Environment solved in 118 episodes! Average Score: 13.06. Time elapsed: 7:33s. See [3](#).

Training DQN 64x32

DQN architecture: fully connected with bias 37x64x32x4. Loss [2](#).

Training history:

Episode 100. Average Score: 4.75. Time elapsed: 3:10




Figure 2: DDQN 32x32: Rewards per episode




Figure 3: DDQN 64x32: Rewards per episode




Figure 4: DQN 64x32: Rewards per episode

Episode 200. Average Score: 11.97. Time elapsed: 6:22

Environment solved in 136 episodes! Average Score: 13.02. Time elapsed: 7:31s. See [4](#).

Training DQN 64x64

DQN architecture: fully connected with bias 37x64x64x4. Loss [2](#).

Training history:

Episode 100. Average Score: 4.75. Time elapsed: 3:21

Episode 200. Average Score: 10.90. Time elapsed: 6:39

Environment solved in 130 episodes! Average Score: 13.03. Time elapsed: 7:38s. See [5](#).

Conclusion

From the above results we can summarize, that DDQN has slightly better training performance. Also 32x32 hidden layer configuration is the best. It was trained faster and has few internal parameters.




Figure 5: DQN 64x64: Rewards per episode

Areas of improvement

During evaluation of the trained agent I noticed, that in some cases the agent hangs up. It begins alternating between 2 positions. You may see it at the end of the video I uploaded on [youtube](#). This is caused by the fact, that the agent makes its decision based only on instant state. It does not take into account previous states and actions. Intuitively is clear that the agent needs some memory to store previous actions and states. This could be achieved by [adding recurrent unit to the network](#), like LSTM, GRU or just simple RNN. Another improvement could be implementing [prioritized experience replay](#).